

AQUABLEND

Deakin University — Capstone Project

In Search of Cluster Analysis

A survey and evaluation of clustering techniques for AquaBlend data

Prepared by

Name	ID	Role
Bao Minh Tran	s224236373	Bridge Lead, Optimisation & Analysis / AI

Reviewed by

Name	ID	Role
<i>Full Name</i>	<i>sXXXXXXXXXX</i>	<i>Reviewer / Supervisor</i>

Version Draft v0.1
Status Draft
Date August 1, 2026

Revision history

Version	Date	Author	Summary of changes
v0.1	August 1, 2026	B. M. Tran	Initial document structure.

Abstract

One paragraph covering: the problem this document addresses, the clustering techniques surveyed, how they were evaluated, and the headline finding or recommendation for AquaBlend.

[TODO: write once the first technique is complete]

Contents

Revision history	i
Notation and acronyms	vi
1 Introduction	1
1.1 Motivation	1
1.2 Objectives	1
1.3 Scope and limitations	1
1.4 Document structure	2
2 Background	3
2.1 Cluster analysis in brief	3
2.2 Families of clustering techniques	4
2.3 Related work	4
3 Data	6
3.1 Source and provenance	6
3.2 Features	6
3.3 Preprocessing	6
3.4 Exploratory summary	6
4 Evaluation protocol	7
4.1 Experimental setup	7
4.2 Validation metrics	7
4.3 Choosing the number of clusters	7
4.4 Interpreting and naming clusters	7
4.5 Threats to validity	7
5 Exploratory Data Analysis and Data Mining Pipeline	8
5.1 Exploratory Data Analysis (EDA) methods	8
5.2 Data Mining pipeline	9
6 Discussion on Dimensionality problems	10
6.1 Curse of Dimensionality	10
6.2 Manifold hypothesis and Manifold learning	10
6.3 Linear Dimensionality Reduction techniques	10
6.3.1 <NAME>	10
6.4 Nonlinear Dimensionality Reduction techniques	11

7	Clustering methods	12
7.1	Dissimilarity measures	12
7.1.1	<NAME>	12
7.2	Hierarchical clustering methods	14
7.2.1	<NAME>	14
7.3	Partition-based clustering methods	15
7.4	Hybrid clustering methods	16
8	Comparative discussion	17
8.1	Side-by-side results	17
8.2	Qualitative comparison	17
8.3	Recommendation	17
9	Conclusion and future work	18
9.1	Conclusion	18
9.2	Future work	18
A	Reproducibility	19
A.1	Environment	19
A.2	Running the analysis	19
B	Supplementary results	19
	Acknowledgements	20
	References	21

List of Figures

List of Tables

1	Validation metrics reported for every technique.	7
2	Validation metrics across all techniques surveyed.	17

Notation and acronyms

Acronyms

Acronym	Expansion
AI	Artificial Intelligence
UML	Unsupervised Machine Learning
HCA	Hierarchical Cluster Analysis
DHC	Divisive Hierarchical Clustering
AHC	Agglomerative Hierarchical Clustering
NC	Number of Cluster
DR	Dimensionality Reduction
EDA	Exploratory Data Analysis
SOM	Self-Organising Map
GNN	Graph Neural Network
DBSCAN	Density-Based Spatial Clustering of Applications with Noise
SSE	Sum of Squared Error

Mathematical notation

Symbol	Meaning
$\mathbf{X}$	Matrix representation of the dataset, rows for samples and columns for features
m	Number of observations
n	Number of features
$\mathcal{C}$	Set of clusters
$\mathcal{C}_i$	i^{th} -cluster of $\mathcal{C}$
$ \mathcal{C} $	Number of clusters
$d(\cdot, \cdot)$	Distance / dissimilarity measure

Conventions

State document-wide conventions here, e.g. vectors are column vectors and set in bold, indices i range over observations and j over clusters, all distances are Euclidean unless stated otherwise, software versions and random seeds used for every run.

1 Introduction

1.1 Motivation

Arising from the desire of improving the project proposal, allowing the core mathematical models to become more proactive, rather than being reactive by nature. In addition, realising the potential of the AquaBlend dataset can be further explored and studied to improve the performance of future Mixed-Integer Nonlinear Programming (MINLP) model, along with insightful patterns undiscovered in the dataset, which is impactful in future research on related problems about water systems based in Melbourne, and possibly adopted to study water systems in other locations.

This document is made to accomodate the needs of summarizing the deliverables of the codebase provided in the project's repository and solid theoretical and practical foundations for exploring and analysing AquaBlend dataset, to facilitate a broad range of functionalities of the project deliverables, including but not limited to, Time-series analysis, Demand Forecasting, Anomaly Detection and Scenarios Generation.

1.2 Objectives

This section outlines questions addressed in this document.

1. What potential clustering methods and dimensionality reduction techniques can be applied to specifically study the AquaBlend dataset?
2. How well the selected methods perform on the AquaBlend dataset subject to certain standard evaluation metrics/scores?
3. Can there be an appropriate Exploratory Data Analysis and Data Mining pipeline to discover patterns in the AquaBlend dataset? If so, can it be standardised to apply to related problems? Otherwise, why is it impossible?
4. Is “Manifold Hypothesis” valid in the dataset? If so, what can be done to improve the clustering results? Otherwise, is there strong statistical evidence attained from the AquaBlend dataset?

1.3 Scope and limitations

Cluster Analysis is a broad research area, thereby we limit ourselves to discussing on specific methods and problems associated, mentioning other methods without going in details. However, this document will be extended to include more clustering methods in the upcoming work.

For the current version, we limit ourselves to those methods and techniques particularly useful to uncover knowledge from AquaBlend dataset. In recent future, we shall improve the pipeline to become adaptable to discover knowledge from relevant datasets aiming to optimise Operational Water Resources Management sector.

More importantly, this document is not a comprehensive survey of all existing studies of Cluster Analysis, and thereby, there may exist more effective and appropriate techniques yet to mentioned within this document. Hence, this document can only be used as a reference for a collection of potential methods, and how they can be practically implemented to solve certain problems.

1.4 Document structure

The document is organised as follows.

1. Survey recent development in Cluster Analysis, see **Sect. 2**.
2. Introduce dataset used (see **Sect. 3**) and experimental setup to benchmark different clustering strategies (see **Sect. 4** and **App. A**).
3. Deliver the Exploratory Data Analysis process and Data Mining pipeline conducted, see **Sect. 5**.
4. Discuss on various problems associated with data dimensionality and possible solutions, see **Sect. 6**.
5. Summarise the theory behind the implementation of clustering methods, see **Sect. 7**.

2 Background

2.1 Cluster analysis in brief

Clustering analysis refers to the process of exploring the dataset to uncover the latent structure and underlying relationships among data points, commonly known as unsupervised machine learning (UML) in the AI community, and essentially important in data mining and knowledge discovery [18]. It comprises of all clustering methods that are capable of handling (usually) high-dimensional data and grouping them by the utilisation of several (usually) quantitative measurements metrics, see [10] for a brief history of development. In this document, we shall refer to a clustering method as [22]:

- Minimising within-cluster (intra-cluster) distance/dissimilarity.
- Maximising intra-cluster distance/dissimilarity.
- Measurement of distance/dissimilarity must be well-defined and practically meaningful.

Formally, let us claim the definitions of a valid mathematical metric and a clustering method.

Definition 2.1 (A mathematical metric [17]). Given two points $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$, a (mathematical) metric is given by $d : \mathbb{R}^n \times \mathbb{R}^n \mapsto \mathbb{R}$, which satisfies:

- Identity and Positivity: $d(\mathbf{x}, \mathbf{y}) > 0$ if $\mathbf{x} \neq \mathbf{y}$, and $d(\mathbf{x}, \mathbf{y}) = 0 \Leftrightarrow \mathbf{x} = \mathbf{y}$
- Symmetry: $d(\mathbf{x}, \mathbf{y}) = d(\mathbf{y}, \mathbf{x})$
- Triangle inequality: $d(\mathbf{x}, \mathbf{y}) \leq d(\mathbf{x}, \mathbf{z}) + d(\mathbf{z}, \mathbf{y}), \forall \mathbf{z} \in \mathbb{R}^n$

Definition 2.2 (Clustering method [18]). Given a dataset $\mathbf{X} \triangleq \{\mathbf{x}_i \in \mathbb{R}^n\}_{i=1}^m$. A clustering method is defined as $\phi_d : \mathbf{X} \mapsto \mathcal{C}$, such that:

$$\bigcap_{i=1}^{|\mathcal{C}|} \mathcal{C}_i = \emptyset \quad \text{and} \quad \bigcup_{i=1}^{|\mathcal{C}|} \mathcal{C}_i = \mathbf{X} \quad (1)$$

and ϕ_d is (usually) aiming to:

$$\min \sum_{i=1}^{|\mathcal{C}|} \sum_{\substack{\mathbf{x}_j, \mathbf{x}_k \in \mathcal{C}_i \\ \mathbf{x}_j \neq \mathbf{x}_k}} d(\mathbf{x}_j, \mathbf{x}_k) \quad \text{and} \quad \max \min_{1 \leq i < j \leq |\mathcal{C}|} \left\{ \min_{\substack{\mathbf{x}_i \in \mathcal{C}_i \\ \mathbf{x}_j \in \mathcal{C}_j}} d(\mathbf{x}_i, \mathbf{x}_j) \right\} \quad (2)$$

where $d(\cdot, \cdot)$ is a *dissimilarity measure*, need not to be a metric as in **Def. 2.1**.

Although different literatures proposing different claims of standard process of cluster analysis [10], though sharing a lot in common, we shall adopt the process written in [22].

1. Feature extraction and selection: opt out representative features of the dataset.
2. Clustering algorithm design: select appropriate methods for the data characteristics and problems' constraints.
3. Cluster evaluation/validation: quantify the performance of clustering results by use of validity measures.

4. Result interpretation: explain what drives the method to identify the patterns.

Looking from the above definition and formulation of clustering methods, several issues have been identified, which are (mostly) subject to dataset and nature of clustering methods [18]. Regarding dataset, the Curse of Dimensionality leads to indistinguishable distances between points in high-dimensional space [15], which degrades the capability of dissimilarity measure used in the methods. Other hurdles are, including but not limited to, size of dataset, noise/outliers, and mixed data type. Regarding methods, many of them requires hyperparameter tuning, which is subjective and data-dependent, also but not limited to, computation complexity, convergence speed, and local optima convergence. Therefore, the upcoming clustering techniques have been developed to mitigate some of issues associated with the problems at some extent.

2.2 Families of clustering techniques

This section serves as a brief list of methods' names, for details what each family include, refer to **Sect. 7**. There is no universal agreement on how clustering methods are classified, several articles, e.g. [21] [6] [10] and books, e.g. [19] and [9] rely on different criteria to divide these clustering methods into groups depending on strategies implemented. In this document, we adopt the philosophy in [18] to classify the methods into two major groups.

Hierarchical clustering Hierarchical Cluster Analysis (HCA) offers multi-level partitions of the dataset, usually represented using a dendrogram, and one can select a threshold to obtain different clustering outcomes. It is further categorised into two groups, i.e. a top-down strategy, known as Divisive Hierarchical Clustering (DHC), and a bottom-up strategy, known as Agglomerative Hierarchical Clustering (AHC). Although differing in construction approaches, they share same components, including dissimilarity measure and linkage criterion. However, DHC is more favorable in practices due to its efficiency.

Partition-based clustering Partitional clustering or iterative relocation algorithms focus on diving the dataset into (nearly-)homogenous groups without the needs of hierarchical structure, aiming to optimise certain criteria functions, can be understood as those given in **Def. 2.2**. It is further categorised into two groups, i.e. hard/crisp clustering and soft/fuzzy clustering. In lieu of considering crisp and fuzzy clustering as two primary groups within partitional methods as in [18] [6], we shall classify crisp clustering methods in more thorough manners and keep fuzzy clustering method as is. Hence, the family of partition-based methods, including but not limited to, graph-theoretic, subspace, density-based, model-based, search-based, SSE-optimised, and miscellaneous clustering methods.

Hybrid clustering this family of method comprises those techniques that are mixture of the other techniques, irrespective of what family they reside in.

In reality, clustering can have any arbitrary shape and distribution, this fact underpins distinct classification of the clustering methods. In brief, there are certain types of clusters that have been extensively studied, e.g. well-separated, prototype-based, graph-based, density-based, and shared-property. For detailed visualisations of these clusters in 2D-space, see [19].

2.3 Related work

Cluster analysis is a fundamental phase in knowledge discovery process, and thus, can be applied to a broad range of real-world applications, e.g. bioinformatics [21], health research [5], and

network-related problems [7]. Nonetheless, we limit ourselves to revisit recent work tackling issues and demands associated with water systems in the past few years.

Two applications of clustering methods, i.e. K-Means and Gaussian mixture clustering, were leveraged by [4], which assisted them to detect recurring operational states, or “fingerprints”, using data obtained from wastewater treatment plants based in Ontario, Canada. Thereafter, successfully identified the associations between these fingerprints and critical events, which allows engineers to identify early signals of critical events from operational conditions. A more closely related research was conducted in Finland, which attempted to model changes in water quality through different stages in a treatment process [11]. They first projected the features onto a lower-dimensional space by the utilisation of Self-organising Map (SOM), and subsequently applied K-Means to cluster the projected data, and successfully identified different prototypes in the data, which implies the variations of water quality through distinct phases conducted in a treatment plant. Indeed, the application of SOM have been widely adopted in water-related research, with two surveys devoted to it in recent years. In particular, [1] systematically reviewed how SOM had been applied to water quality assessment in reservoirs and lakes. Meanwhile, [14] highlighted the improved interpretability of SOM when being merged with a clustering algorithm, as studied in [11], and superiority over other relevant methods in modelling the nonlinear relationship between variables, and can be well visualised in 2D or 3D space.

Aside from the scope of water quality, the majority of the literature focusing on discovery of water demands. More specifically, A comparative study on daily consumption patterns evaluates *K*-means, agglomerative hierarchical and spectral clustering under the Silhouette Coefficient and the Calinski–Harabasz index [8]. Closer to density-based methods, a modified DBSCAN with automatic optimisation of the neighbourhood radius and explicit detection of noise points has been applied to segment daily consumption profiles [12]. Time-series formulations are also represented, by clustering of consumption series combined with entropy analysis [20] and by clustering of hydrological telemetry series [16]. More recently, a graph-based clustering method, specifically a Graph Neural Network (GNN), was proposed in [3], to leverage the interdependencies between the dataset and external information unavailable to it, which is likely to govern the urban water demands.

A relatively sparse research direction is the problem-driven scenario clustering. Typically, scenarios are generated by sampling from the distributions of the operational dataset, and sometimes overlook the mathematical models, as highlighted in [13]. A recent literature, see [2], not only worked on this problem-driven scenario clustering method, but also emphasised the pivotal role of feature engineering to yield better clustering results. Both of them inspired a more appropriate scenario generation approaches to comply with the real-world cases while considering the mathematical model used. This is essentially important, as the mathematical model will soon be developed to become a stochastic programming model, and in reality, the parameters given to the model are not fixed.

Up to our knowledge, no recent work connects regimes discovered in an operational water-quality feature space to source-blending or source-selection decisions, and we suppose this is a potential and abundant area to discover and study.

3 Data

3.1 Source and provenance

Where the data comes from, who owns it, which extract or version is used, and the date it was pulled.

3.2 Features

Table of features: name, type, unit, range, and whether the feature is used for clustering or held out for interpretation.

3.3 Preprocessing

Cleaning, missing-value handling, encoding, scaling, and dimensionality reduction — everything applied before clustering, in the order applied, so the pipeline is reproducible.

3.4 Exploratory summary

Distributions, correlations, and anything that constrains the choice of technique (cluster shape, density, outliers, scale).

4 Evaluation protocol

4.1 Experimental setup

Hardware, software versions, random seeds, and the number of restarts per run. Identical for every technique.

4.2 Validation metrics

The internal and external validity measures reported for every technique, with the definition of each given once here.

Table 1: Validation metrics reported for every technique.

Metric	Range	Interpretation
<i>Metric</i>	<i>[a, b]</i>	<i>What a high or low value means</i>

4.3 Choosing the number of clusters

The procedure applied consistently wherever k must be fixed in advance.

4.4 Interpreting and naming clusters

How a resulting partition is turned into something the business can act on, and who signs off on the naming.

4.5 Threats to validity

Known weaknesses of this protocol and how they are mitigated.

5 Exploratory Data Analysis and Data Mining Pipeline

5.1 Exploratory Data Analysis (EDA) methods

5.2 Data Mining pipeline

6 Discussion on Dimensionality problems

6.1 Curse of Dimensionality

6.2 Manifold hypothesis and Manifold learning

6.3 Linear Dimensionality Reduction techniques

6.3.1 <NAME>

Overview *What the technique does, in one paragraph: what it maps from, what it maps to, and what it tries to preserve in doing so. State whether it is linear or nonlinear, and for a nonlinear technique, its relation to the manifold hypothesis of Section 6.*

Assumptions and applicability *What the technique assumes about the data — linearity, local versus global structure, density, scale, noise — and the conditions under which it is a sensible choice for the data described in Section 3. State explicitly whether it is intended for visualisation, for reducing input to a clustering method, or for both: these are different jobs and a technique good at one may distort what the other needs.*

Formulation *The objective being optimised and the mapping obtained, using the symbols declared in the notation table. Cite the original source.*

Out-of-sample and inverse mapping *Two questions, answered explicitly, because the rest of the pipeline depends on both.*

First: is the technique inductive or transductive? That is, does it learn a mapping that applies to observations unseen during fitting, or does it embed only the sample it was given? A transductive technique cannot precede a clustering step in a pipeline that will later see new data.

Second: does an inverse mapping exist? If so, a component can be read back in terms of the measured features; if not, say so, and say where the interpretation of the reduced space will come from instead.

Hyperparameters and tuning *Each hyperparameter, the range searched, and the value selected — with the reason. For neighbourhood-based techniques, state what the neighbourhood size does to the balance between local and global structure, since it is the parameter most likely to change the conclusions drawn from a figure.*

Choosing the number of components *How the target dimension was chosen, and on what evidence: an explained-variance criterion, an intrinsic dimension estimate (see Section 6), a reconstruction error, or the requirements of a downstream method. Choosing two because two is what plots is a legitimate choice for a figure and not for an analysis — if that is the reason, say so.*

Complexity *Time and space complexity, and what that means at AquaBlend data volumes. Note any step that is quadratic in the number of observations, since that is what decides whether the technique is usable on the full dataset or only on a sample.*

Application to AquaBlend data *Implementation: library and version, random seeds, and how the run is reproduced. Then any deviation from the shared preprocessing pipeline of Section 3.3, and why it was necessary — scaling in particular, since most of these techniques are not scale-invariant.*

Results *Two things, in this order.*

Structure preservation: how faithfully the embedding represents the input, reported quantitatively — explained variance, trustworthiness, stress, or reconstruction error, whichever the technique admits — and not by appeal to the appearance of a scatter plot.

Effect on downstream clustering: the validity measures of Section 4.2 for the same clustering method fitted with and without this reduction. This is the part that decides whether the technique earns its place in the pipeline.

Limitations *Where this technique distorted the data, or would be expected to. Distinguish what was observed on our data from what is known from the literature. For nonlinear techniques, state plainly which properties of the embedding must not be read as properties of the data — typically inter-cluster distances, relative cluster sizes, and apparent density.*

Summary *Two or three sentences: the verdict on this technique, and whether it is recommended for visualisation, for the clustering pipeline, for both, or for neither. Phrased so it can be carried directly into Section 8.*

6.4 Nonlinear Dimensionality Reduction techniques

7 Clustering methods

Refer to **Sect. 2** for details of cluster analysis, families of methods, and important properties.

7.1 Dissimilarity measures

Add a short paragraph distinguishing between dissimilarity measure vs. validity index

7.1.1 <NAME>

Overview *What the measure quantifies, in one paragraph, and which kind it is: a dissimilarity measure, or a validity index — and if an index, whether internal, external or relative. State why it is included here: which methods need it, or which question about a partition it answers.*

Definition *The formal definition, using the symbols declared in the notation table. Cite the original source. Where the measure is defined only for particular data types or particular cluster representations, state that here rather than in a later part.*

Properties *For a dissimilarity measure: which of the three properties of Definition 2.1 it satisfies — identity and positivity, symmetry, the triangle inequality — and which it violates. This is not a formality. A measure that violates the triangle inequality is admissible under Definition 2.2, but cannot be used with a method whose correctness depends on it.*

For a validity index: the attainable range, the direction (whether high or low is better), whether it is corrected for chance, and whether it is defined when observations are labelled as noise.

Applicability *Which data types the measure accepts — numeric, categorical, mixed, incomplete — and which families of Section 2.2 it can serve.*

For an index, add: whether it applies to a soft partition directly or only to a defuzzified one, and how observations labelled as noise are treated. Excluding them scores only the points a method was confident about, which flatters that method; whichever choice is made, it must be stated wherever the index is reported.

Computation and complexity *How the measure is computed, its time and space complexity, and what that means at AquaBlend data volumes. Note in particular whether it requires the full pairwise matrix, since that is quadratic in the number of observations and often the binding constraint.*

Application to AquaBlend data *Implementation, and where the measure is used: which methods are fitted with it, or which results are reported under it. Any parameter it carries, with the value chosen and the reason.*

Behaviour *What the measure rewards, demonstrated rather than asserted — on our data, on a benchmark with known structure, or on a synthetic example.*

For an index this is essential, because every index encodes a notion of what a cluster is. An index built on distances to a centroid rewards compact, isotropic clusters and penalises elongated or

irregular ones, so it does not rank methods of different families neutrally. Where that applies, say so here and carry it into Section 4.5.

Limitations *Where the measure is uninformative, misleading, or undefined — degenerate partitions, unbalanced cluster sizes, high dimension, or the concentration of distances discussed in Section 6. Distinguish what was observed from what is known from the literature.*

Summary *Two or three sentences: what this measure is good for, what it must not be used for, and whether it is part of the reported protocol of Section 4.2 or supplementary.*

7.2 Hierarchical clustering methods

7.2.1 <NAME>

Overview *What the algorithm does, in one paragraph, and the family it belongs to (see Section 2.2).*

Assumptions and applicability *What the algorithm assumes about cluster shape, density, scale and noise; the conditions under which it is a sensible choice for the data described in Section 3.*

Formulation *Objective function and update rule, using the symbols declared in the notation table. Cite the original source.*

Hyperparameters and tuning *Each hyperparameter, the range searched, and the value selected — with the reason.*

Complexity *Time and space complexity, and what that means at AquaBlend data volumes.*

Application to AquaBlend data *Implementation, any deviations from the shared preprocessing pipeline, and why they were necessary.*

Results *Metrics from Section 4.2, cluster profiles, and supporting figures.*

Limitations *Where this method failed or would be expected to fail on this data.*

Summary *Two or three sentences: the verdict on this method, carried into the comparison section.*

7.3 Partition-based clustering methods

7.4 Hybrid clustering methods

8 Comparative discussion

8.1 Side-by-side results

Table 2: Validation metrics across all techniques surveyed.

Technique	Metric 1	Metric 2	Runtime
Technique one	—	—	—

8.2 Qualitative comparison

Trade-offs the metrics do not capture: interpretability, stability across seeds, sensitivity to preprocessing, operational cost.

8.3 Recommendation

Which technique AquaBlend should adopt, under what conditions, and what would change the recommendation.

9 Conclusion and future work

9.1 Conclusion

The answer to each objective listed in Section 1.2, in the same order.

9.2 Future work

Techniques not yet covered, open questions, and what the next contributor to this document should pick up first.

A Reproducibility

A.1 Environment

Language, library and package versions, and how to recreate the environment.

A.2 Running the analysis

Commands to reproduce every result in this document, and where the code and outputs live in the repository.

B Supplementary results

Full metric tables, additional figures, and parameter sweeps that would interrupt the main narrative.

Acknowledgements

Thank the people and bodies that contributed: teammates who supplied or cleaned the data, reviewers, supervisors and unit staff, and any external dataset or software provider whose licence requires attribution. Name the contribution, not just the person.

Note here any use of generative AI tools, stating what they were used for and how the output was verified.

References

- [1] Oraib Almegdadi, João Marcelino, Sarah Fakhreddine, João Manso, and Nuno C. Marques. Self-organizing maps for water quality assessment in reservoirs and lakes: A systematic literature review, 2026.
- [2] Majdi Argoubi and Khaled Mili. Stochastic multi-objective optimization for flood control in multi-reservoir systems: an adaptive progressive hedging approach with scenario clustering. *Frontiers in Water*, Volume 7 - 2025, 2025.
- [3] Jonnatan Arias-Garcia, David Cárdenas-Peña, Álvaro Angel Orozco-Gutiérrez, Hernán Felipe Garcia-Arias, and Jhoniers Gilberto Guerrero-Erazo. Graph-based clustering of urban water consumption profiles via adaptive attention and multi-relational topologies. *Water*, 18(11), 2026.
- [4] Michael De Santi, Ahmed Al-Sayed, Stephanie Gora, Ciprian Panfilie, Sangeeta Chopra, Jianping Zhang, Xiaohui Jin, Dale Barker, Yaldah Azimi, Satinder Brar, and Usman Khan. Exploring wastewater treatment plant operation and performance using k-means and gaussian mixture clustering. *Water Science and Technology*, 92(2):251–268, 07 2025.
- [5] Bei Dou, Bert Vandenberk, and Philip Moons. Cluster analysis: An overview of traditional and novel approaches for health researchers. *European Journal of Cardiovascular Nursing*, 2026. Advance online publication, article zvag100.
- [6] Absalom E. Ezugwu, Abiodun M. Ikotun, Olaide O. Oyelade, Laith Abualigah, Jeffery O. Agushaka, Christopher I. Eke, and Andronicus A. Akinyelu. A comprehensive survey of clustering algorithms: State-of-the-art machine learning applications, taxonomy, challenges, and future research prospects. *Engineering Applications of Artificial Intelligence*, 110:104743, 2022.
- [7] Ali Noori Gatea, Haider Sh. Hashim, Hamid Ali Abed AL-Asadi, Didem Kıvanç Türeli, Zaid Ameen Abduljabbar, and Vincent Omollo Nyangaresi. Clustering in data analysis: Comprehensive insights into techniques and challenges. In Radek Silhavy and Petr Silhavy, editors, *Software Engineering: Emerging Trends and Practices in System Development*, pages 108–124, Cham, 2026. Springer Nature Switzerland.
- [8] Hongyuan Guo, Xingpo Liu, and Qichen Zhang. Identifying daily water consumption patterns based on k-means clustering, agglomerative hierarchical clustering, and spectral clustering algorithms. *AQUA - Water Infrastructure, Ecosystems and Society*, 73(5):870–887, 04 2024.
- [9] Jiawei Han, Micheline Kamber, and Jian Pei. *Cluster Analysis: Basic Concepts and Methods*, chapter 10. The Morgan Kaufmann Series in Data Management Systems. Morgan Kaufmann, Waltham, MA, 3 edition, 2011.
- [10] Adam Jaeger and David Banks. Cluster analysis: A modern statistical review. *WIREs Computational Statistics*, 15(3):e1597, 2023.
- [11] P. Juntunen, M. Liukkonen, M. Lehtola, and Y. Hiltunen. Cluster analysis by self-organizing maps: An application to the modelling of water quality in a treatment process. *Applied Soft Computing*, 13(7):3191–3196, 2013.
- [12] Aliaksey A. Kapanski, Roman V. Klyuev, Vladimir S. Brigida, and Nadezeya V. Hruntovich. Temporal segmentation of urban water consumption patterns based on non-parametric density clustering. *Technologies*, 13(10), 2025.

- [13] Julien Keutchayan, Janosch Ortmann, and Walter Rei. Problem-driven scenario clustering in stochastic optimization. *Computational Management Science*, 20(1):1–33, December 2023.
- [14] Sabina Licen, Aleksander Astel, and Stefan Tsakovski. Self-organizing map algorithm for assessing spatial and temporal patterns of pollutants in environmental compartments: A review. *Science of The Total Environment*, 878:163084, 2023.
- [15] Dehua Peng, Zhipeng Gui, and Huayi Wu. Interpreting the curse of dimensionality from distance concentration and manifold effect, 2025.
- [16] Intouch Prakaisak and Papis Wongchaisuwat. Hydrological time series clustering: A case study of telemetry stations in thailand. *Water*, 14(13), 2022.
- [17] Walter Rudin. *Principles of Mathematical Analysis*. International Series in Pure and Applied Mathematics. McGraw-Hill, New York, 3 edition, 1976.
- [18] Jaswinder Singh and Damanpreet Singh. A comprehensive review of clustering techniques in artificial intelligence for knowledge discovery: Taxonomy, challenges, applications and future prospects. *Advanced Engineering Informatics*, 62:102799, 2024.
- [19] Pang-Ning Tan, Michael Steinbach, Anuj Karpatne, and Vipin Kumar. *Cluster Analysis: Basic Concepts and Algorithms*, chapter 7. Pearson Education, New York, NY, 2 edition, 2019.
- [20] Renfang Wang, Xinyu Zhao, Hong Qiu, Xu Cheng, and Xiufeng Liu. Uncovering urban water consumption patterns through time series clustering and entropy analysis. *Water Research*, 262:122085, 2024.
- [21] Aasim Ayaz Wani. Comprehensive analysis of clustering algorithms: exploring limitations and innovative solutions. *PeerJ Computer Science*, 10:e2286, 08 2024.
- [22] Dongkuan Xu and Ying jie Tian. A comprehensive survey of clustering algorithms. *Annals of Data Science*, 2:165 – 193, 2015.